



OUTLOUD V2 // AI PM CASE STUDY

OutLoud V2

A judgment-free voice AI coach for Indian job seekers, now for interviews and the workplace too

From user research to a live, tested app: 3 AI modes, 8 build phases with an AI coding agent, 63 logged tests

Nisha Verma (Megha) · Product Manager · October 2026



They know English. They freeze when they have to speak it.

The gap isn't knowledge. It's speaking live, unscripted, under pressure.

Early-career workers and job seekers in India read and write English, yet hesitate in interviews and at work.

Coaching means strangers and fixed schedules; apps paywall or drill grammar. There is no private, everyday way to practise speaking.

90.9%

would use a free app to practise
speaking privately every day

72.8%

say weak spoken English has cost,
or could cost, them an opportunity

27.3%

name fear of being judged as a top
blocker to daily practice

Meet Sneha, 26 (Ranchi): the capability–confidence gap

Same persona as V1. In V2, each of her constraints became a design rule.

What she can do

Reads and writes English accurately
Teaches English–medium academic content
Has deep subject knowledge and comprehension

Where she gets stuck

Freezes in interviews and unscripted talk; it has cost her jobs
No one around her speaks English; her coaching support is gone
Won't pay for another paywalled app, and fears being judged

What V2 does about it

"Free to use. No payment, ever." on Login and Sign Up
Interview and Workplace modes for her real career moments
Never corrects accent, fillers or Indian English

What users told me decided what V2 would add.

Their own words → the pattern behind them → what V2 does about it

In their words	Insight	V2 implication
<i>"But my mind just goes blank when I need to open my mouth."</i> Rony, ex-barista	Judgment stops people starting: 27.3% name it a top blocker.	No scores, no accent correction; max 2 feedback points.
<i>"...if I try teaching in English, I just lose the essence and students do not understand."</i> Rashmi, TET-certified teacher	The stakes are career moments: 40.9% are preparing for exams or interviews.	Interview and Workplace modes, shaped by her target role.
<i>"I don't have enough time to spend on coaching first... I want to do practice while doing my work."</i> Bhavya, restarting her career	Time is why past attempts failed: 36.4% quit for lack of it.	3 minutes counts as a completed session; reminders on her schedule.

V1 proved people would speak. It didn't prove they'd come back.

Live V1 data: the product worked, but the practice habit didn't form.

58
sessions, 28 devices

22.4%
completion rate

14 / 16
said "Yes, helpful"

2.3 min
average session

What V1 showed	What V2 does about it
Users came for interviews and work	Interview and Workplace modes
Feedback was a top retention driver	Spoken feedback on demand and at the end
No way to recognise returning users	Accounts, verification and a profile
Reminders: top retention lever (27.3%)	User-set email reminders

General AI chat can talk. It can't coach you for Monday's interview.

Some target users already practise with ChatGPT and Claude.

General AI chat apps

- Built to answer, not to coach
- You have to prompt it into a coach
- Corrects whatever it notices
- Doesn't know your role or interview
- Nothing brings you back to practise

VS

OutLoud V2

- ✓ Voice-first: speak, listen, reply
- ✓ Three modes for real situations
- ✓ Never corrects accent or Indian English
- ✓ Knows your profession and target role
- ✓ Reminders on your own schedule

Positioning: free, private, voice-first practice for the moments that decide a career.

I scoped V2 around the question V1 left open.

V1 asked: will users speak? V2 asks: can practice become personal and habitual?

① Built in V2

Personal + habitual

Three modes, three prompts
Accounts + email verification
Profile reused in prompts
Spoken feedback, 2 modes
Email reminders she schedules
Data isolation on the server

② Deferred

Worth it, not yet

Streaming speech and voice
Exact-words speech-to-text
Guilt-free re-entry test
Talk-ratio metrics
Paid AI, only if usage grows

③ Cut

Out of V2 scope

Saved history: cost grows
Progress view
Admin notifications
Terms of Service step
V1 data linking

Scope test, carried over from V1: does this strengthen the core problem? If not, it waited or it was cut.

Impact × Effort: why some features shipped and others waited

Impact = effect on Week-4 completion. Effort = build time + free-plan running cost.

LOW EFFORT →

→ HIGH EFFORT

HIGH
IMPACT ↑

Quick wins

BUILT FIRST

Email reminders (free email service)
Profile asked once, reused in every prompt
Feedback through hidden button signals

Big bets

BUILT / DEFERRED

Three modes, three prompts: built
Accounts + verification: built
Streaming voice: deferred

LOW
IMPACT

Fill-ins

DROPPED

Terms of Service step
Admin-pushed notifications

Money pits

CUT

Saved conversation history
Progress view
Cross-device V1 data linking

One coach can't serve three situations, so V2 has three modes.

Each mode has its own system prompt, and correction works differently in each.

General

Relaxed everyday conversation to build comfort speaking
Gentle inline correction, max one per turn

Interview

A realistic mock interview, adapted to the user's target role
No mid-interview correction;
feedback on request and at the end

Workplace


AI asks what to practise (a meeting, a pitch, a client call), then role-plays it
Selective correction plus session feedback

Personalised from a one-time profile — profession, career goal, target role — so Interview and Workplace feel relevant from the first session.

General builds on the tested V1 prompt; Interview and Workplace have their own prompts

Every screen removes friction from the next behaviour.

Profile once → choose a mode → speak → feedback → a reminder brings you back



Design principle: OPEN → SPEAK → COMPLETE → RETURN · 10 screens · no back or menu icon inside a session

I reversed my own V1 decision: an account before the first word.

V1 asked only for a name. V2 asks for an account and a short profile, on purpose.

V1: name only

- Fastest possible start: one field
- Anonymous device ID in the browser
- No reliable returning-user view
- Couldn't personalise practice

V2: account + profile

- Email, password and verification
- Name, profession, career goal, target role
- The same user on any device
- Interview questions fit her target role

The cost, and how I reduced it

- Friction** before the first session → "Free to use. No payment, ever." shown on Login and Sign Up
- Profile** asked once, editable any time in Settings
- V1 history** can't be linked: accepted

Why: no accounts means no Week-4 metric, no reminders and no personalisation.

I broke one V1 rule on purpose: feedback at the end.

V1 had no end-of-session report. Interview and Workplace users need a clear signal, so V2 adds one.

Mode	During the session	Get Feedback button	When the session ends
General	Gentle inline correction, max 1 per turn	Not shown	No report, same as V1
Interview	None, so it feels like a real interview	Up to 2 spoken points	Up to 2 spoken points
Workplace	Selective inline correction	Up to 2 spoken points	Up to 2 spoken points

- Spoken, never a scorecard or a grade
- Each point tied to something she actually said
- Never accent, fillers or Indian English
- Never repeated; none if she hasn't said enough

The user decides when we ask her to come back.

Coming back is still the hardest moment. Reminders (top lever, 27.3%) had to help without nagging.

V2 choice	Why	Rejected
She sets day, time and timezone	She controls it, so it doesn't feel like nagging	App-chosen schedules
Email, not push	Works on every phone; web push needs a home-screen install on iPhone	Web push
No streaks or missed days	Guilt is what stops people coming back (V1 insight)	Streak counters
Timer wakes the app only when due	A sleeping free-plan server would never send it otherwise	Always-on server

Tested end to end: reminder delivered by Brevo's free email service in about 4 seconds, twice.

Product principles became rules the AI can't break.

Three system prompts, written before any code and loaded word for word.

Never, in every mode

Correct an accent, fillers, pauses or Indian English patterns

Give scores, grades or a list of mistakes

Ask more than one question at a time

End the session for the user

Invent details about the user's experience

Hidden control signals

Buttons send [GET_INTERVIEW_FEEDBACK] or [FINISH_WORKPLACE]; the AI acts on them and never says them aloud

Evidence-based feedback

Max two points, each tied to something the user actually said; never repeated; none if there isn't enough yet

Speech-aware reading

Workplace reads spoken structure ("first... second...") from unpunctuated transcripts

Minimal context

General gets only the name; Interview and Workplace get the profile

I tested what the AI does, not just what the prompt says.


Some behaviour passed, some was fixed, some I chose to accept. All of it is logged.

Observed	Decision	Result
Voice silently cut off replies over ~200 characters	Split replies into chunks of 180 characters or fewer	Fixed: found by comparing audio length to text
Hidden button signals could be spoken aloud	Prompt rule: act on signals, never say them	Passed: 3 of 3 leak tests clean
General often opens with the same question	Accepted; prompt left unchanged	Logged as a known behaviour
Speech-to-text tidies small grammar slips	Accepted for V2	Biggest limit on correction: next bet


The AI never decides the session is over. The user does.

Locked rule: user says bye → AI finishes → mic off → wait for End Session.

User says bye




User taps End Session first (Interview, Workplace)



A session counts as **completed** at 3 minutes of practice or a natural bye. Feedback requests are never mistaken for a bye — a rule I specified and the tests confirmed.

A full voice AI stack on a ₹0 budget, by choice.

Every service runs on a free plan. I measured what that costs users before accepting it.



An AI agent wrote the code. I made every product call.

Claude Code built V2 in 8 phases; nothing moved forward without my review.

Phase 1

App skeleton, live on a phone

Phase 2

Database, V1 data untouched

Phase 3

Accounts, onboarding,
settings

Phase 4

Voice loop + General mode

Phase 5

Interview + Workplace,
feedback

Phase 6

Email reminders

Phase 7

Security hardening

Phase 8

Launch checks

Every phase, the same loop

Agent's plan → my review and decisions → build → automated tests → my phone test → commit → deploy

What review caught

A sleeping-server reminder trap, an always-on deployment eating credits, and an unneeded database table

The worst bug threw no error. A test found it anyway.

63 logged tests across 8 phases, with evidence, including the ones only accepted.

63

logged tests, 60 passed

40/41

launch checks passed

8

bugs caught before release

0

secrets in code or git

What a test caught	Fix
Voice silently cut off replies over ~200 characters	Replies split into sentence chunks
A V1 database rule blocked every V2 session	Column made optional; V1 data untouched
App crashed on deploy (wrong Node version)	Node 22, and no crash on bad settings
Security: logged-out and cross-user access	All attempts refused; rate limits exact


I made trade-offs explicit before they became product problems.

Every decision was logged with why it was made and what it cost.

Decision	Why	Trade-off accepted
Free plans only	Portfolio scale, no revenue to fund AI costs	Daily voice limit; text fallback
Email verification	A typo email would break reset and reminders	One extra step before practice
Drop V1 data linking	Different web addresses can't share browser storage	72 V1 rows stay unlinked
App never auto-exits	Slow readers never lose their final feedback	One extra tap on End Session
In-memory rate limits	One app instance, so no new database table	Counts reset when the app sleeps
Transcripts never stored	Privacy, and costs that don't grow with use	No talk-ratio metric in V2

V2 is live. The next job is moving real users through the loop.

Same problem—first outreach as V1, plus the step V1 was missing: bringing users back.



First milestone: 3–5 real users through the whole loop, on iPhone and Android.

I kept the North Star and made it measurable for each mode.

The question stays the same: are users still practising in Week 4?

SESSION DURATION

5–10 min

average per session

V1 actual: 2.3 min

INTERVIEW

≥ 40%

interview practice completion

NORTH STAR METRIC

≥ 40%

Week-4 practice completion

Completed = 3 min of practice or a natural bye.

Active user = starts at least one session that week.

WORKPLACE

≥ 1 / week

completed session per active user

TALK RATIO

70–80%

user speaking time vs AI

Not measurable in V2: no transcripts

My next question isn't 'What feature should I add?'

It's 'What is speech-to-text hiding from the AI?'

HYPOTHESIS

Speech-to-text tidies small grammar slips before the AI sees them, so users get fewer useful corrections than they should.

TEST

In the 3–5 person real-user test, check whether testers notice corrections the AI missed; then trial a setup that keeps exact words.

SUCCESS METRIC

More real mistakes caught per session, and a higher "Yes, helpful" share on the feedback screen.

FINAL PM TAKEAWAY

V1 proved people will speak to an AI coach.

V2 has to prove they come back, and that the coaching is accurate when they do.

Try it: outloud-v-2.replit.app ·

Full case study (58 pages) [<https://rb.gy/Oav5kh>]